



Individuallität. Unisong. Divisio-

Präsumt-Lösung.

Über-

Erinnerung.

Erinner-

Sprache.



**BATCH :** **BATCH 150 DATA SCIENCE**

---

**LESSON :** **DEEP LEARNING**

---

**DATE :** **21.08.2023**

---

**SUBJECT :** **Deep Learning  
and ANN**



-  techproeducation
-  techproeducation
-  techproeducation
-  techproeducation
-  techproedu




# DEEP LEARNING


TECHPRO  
EDUCATION

# Kahoot!


**ANN = 18 saat**




**CNN = 21 saat**



**COMPUTER VISION = 24 saat**




**RNN+LSTM+GRU = 12 saat**



**CAPSTONE  
PROJECT**

- $c_{t-1}$  previous cell state
- $f_t$  forget gate output
- $i_t$  input gate output
- $c_t$  candidate
- $c_t$  new cell state

**NLP = 18 saat**



**TOPLAM  
93 saat**



# DEEP LEARNING




ChatGPT



# LESSON CONTENTS

- 🕒 Why Deep Learning?
- 🕒 AI and Deep Learning Milestones
- 🕒 Artificial Neural Network (ANN)
- 🕒 Deep Neural Network
- 🕒 Activation Functions
- 🕒 Gradient Descent
- 🕒 Backpropagation
- 🕒 Regularization
- 🕒 Most Popular Deep Learning Libraries










# WHY DEEP LEARNING?

---



1. Machine Learning Engineer
2. Deep Learning Engineer
3. Computer Vision Engineer
4. Natural Language Processing Engineer
5. AI Research Scientist
6. AI Software Developer
7. AI Consultant
8. Data Scientist - AI/ML



## How much does a Computer Vision Engineer make in USA?

**\$165,000** / Annual


Based on 968 salaries

The average computer vision engineer salary in the USA is \$165,000 per year or \$84.62 per hour. Entry level positions start at \$137,500 per year while most experienced workers make up to \$205,000 per year.



## Data Scientist Average Salary

ComputerCareers.org



## How much does a Natural Language Processing Engineer make?

Experience

All years of Experience

**\$88,551** / yr

Total Pay

\$75,072 / yr

Base Pay

\$13,479 / yr

Additional Pay

\$88,551 / yr

\$49K

\$28K

\$162K

\$279K

■ Most Likely Range ■ Possible Range



## AI job titles with the highest salaries



Rank	Job title	Average salary
1.	Machine learning engineer	\$142,858.57
2.	Data scientist	\$126,927.41
3.	Computer vision engineer	\$126,399.81




Prestige®



# DEEP LEARNING



## DEEP LEARNING EVERYWHERE



### INTERNET & CLOUD

Image Classification  
Speech Recognition  
Language Translation  
Language Processing  
Sentiment Analysis  
Recommendation

### MEDICINE & BIOLOGY

Cancer Cell Detection  
Diabetic Grading  
Drug Discovery

### MEDIA & ENTERTAINMENT

Video Captioning  
Video Search  
Real Time Translation

### SECURITY & DEFENSE

Face Detection  
Video Surveillance  
Satellite Imagery

### AUTONOMOUS MACHINES

Pedestrian Detection  
Lane Tracking  
Recognize Traffic Sign



# DEEP LEARNING

## 10 FASCINATING APPLICATIONS OF DEEP LEARNING



SELF-DRIVEN  
CARS



AUTOMATIC  
HANDWRITING  
GENERATION



PIXEL  
RESTORATION



COLOURISATION OF  
BLACK & WHITE  
IMAGES



DEEP  
DREAMING



DETECTION OF  
GROWTH DELAYS  
IN CHILDREN



DEMOGRAPHIC  
PREDICTION



SOUND  
ADDITION TO  
SILENT FILMS



NEWS  
GENERATION





AUTOMATIC  
MACHINE  
TRANSLATION

## 20 DEEP LEARNING Applications

- 
- 1 Self Driving Cars
  - 2 Entertainment
  - 3 Visual Recognition
  - 4 Virtual Assistants
  - 5 Fraud Detection
  - 6 Natural Language Processing
  - 7 News Aggregation and Fraud News Detection
  - 8 Detecting Developmental Delay in Children
  - 9 Colourisation of Black and White images
  - 10 Adding sounds to silent movies
  - Healthcare 11
  - Personalisations 12
  - Automatic Machine Translation 13
  - Automatic Handwriting Generation 14
  - Demographic & Election Predictions 15
  - Automatic Game Playing 16
  - Language Translations 17
  - Pixel Restoration 18
  - Photo Descriptions 19
  - Deep Dreaming 20



# AI/MACHINE LEARNING/DEEP LEARNING





# WHAT IS DEEP LEARNING?



## ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior




## MACHINE LEARNING


Ability to learn without explicitly being programmed



## DEEP LEARNING

Extract patterns from data using neural networks







- Artificial intelligence
- Machine learning
- Language Processing
- Deep learning



# AI/MACHINE LEARNING/DEEP LEARNING



Jeff Winter






# DEEP LEARNING



İnsan beyninin çalışmalarını taklit eden yapay zeka işlevidir.



Derin öğrenme bir makine öğrenimi yöntemidir!



Büyük miktarda veri gerektirir.



# WHY DEEP LEARNING IS SO POPULAR

TECHPRO  
EDUCATION

## Why Now?

Neural Networks date back decades, so why the resurgence?

### 1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



WIKIPEDIA  
The Free Encyclopedia



### 2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



### 3. Software

- Improved Techniques
- New Models
- Toolboxes


TensorFlow



**RAPIDLY INCREASING DATA**

---

# Why deep learning



How do data science techniques scale with amount of data?

# WHY DEEP LEARNING IS SO POPULAR



AT'S A ZETTABYTE?

1 kilobyte	1,000,000,000,000,000,000
1 megabyte	1,000,000,000,000,000,000
1 gigabyte	1,000,000,000,000,000,000
1 terabyte	1,000,000,000,000,000,000
1 petabyte	1,000,000,000,000,000,000
1 exabyte	1,000,000,000,000,000,000
1 zettabyte	1,000,000,000,000,000,000

Data volume in zettabytes

150

100

50

0



# 2022

PEOPLE SEND 16M

TEXTS



ONLINE  
EVENT  
GOERS PURCHASE  
\$12.9K

231.4M messages

EMAIL  
USERS SEND

GOOGLE  
USERS CONDUCT

5.9M searches



FACEBOOK

USERS SHARE

1.7M pieces of content



INSTAGRAM  
USERS SHARE

66K photos



TWITTER  
USERS SHARE

347.2K tweets



CRYPTO  
BUYERS PURCHASE

\$90.2M in cryptocurrency



VENMO  
USERS SEND

\$437.6K



AMAZON  
SHOPPERS SPEND

\$443K



EVERY  
MINUTE  
OF THE DAY



SNAPCHAT  
USERS SEND

2.43M snaps



TINDER  
USERS SWIPE

1.1M times



YOUTUBE  
USERS UPLOAD

500 hours of video



VIEWERS SPEND  
1M hours  
STREAMING



DOORDASH  
DINERS PLACE

\$76.4K in orders





104.6K hours  
SPENT IN  
ZOOM  
MEETINGS





# TECHNOLOGICAL ADVANCEMENTS IN HARDWARE

---



# Generations and Future Computers





# WHY DEEP LEARNING IS SO POPULAR

TECHPRO  
EDUCATION

## THE BIG BANG IN DEEP LEARNING



DNN




BIG DATA



GPU

*“The GPU is the workhorse of modern A.I.”*

POPULAR  
SCIENCE



GPU AND Deep Learning




**GPU and Deep Learning:**  
A Combination That Works Miracles




# WHY DEEP LEARNING IS SO POPULAR

TECHPRO  
EDUCATION




## CPU

- Small models
- Small datasets
- Useful for design space exploration



## GPU

- Medium-to-large models, datasets
- Image, video processing
- Application on CUDA or OpenCL



## TPU


- Matrix computations
- Dense vector processing
- No custom TensorFlow operations

Central processing unit



CPU

Graphics processing unit



GPU

Tensor Processing Unit



TPU





# WHY DEEP LEARNING IS SO POPULAR

TECHPRO  
EDUCATION



CPU



GPU



Tensor Processing Unit

TPU



# WHY DEEP LEARNING IS SO POPULAR

TECHPRO  
EDUCATION

## Why GPU Matters in Deep Learning?

```
X_train shape: (50000, 3, 32, 32)
50000 train samples
10000 test samples
Using real-time data augmentation.
Epoch 1/200
50000/50000 [=====] 734s
```

Epoch 2/200  
50000/50000 [=====] 733s  
Epoch 3/200  
50000/50000 [=====] 733s  
Epoch 4/200  
50000/50000 [=====] 733s

**VS**

```
X_train shape: (50000, 3, 32, 32)
50000 train samples
10000 test samples
Using real-time data augmentation.
Epoch 1/200
50000/50000 [=====] 27s
```

Epoch 2/200  
50000/50000 [=====] 27s  
Epoch 3/200  
50000/50000 [=====] 27s  
Epoch 4/200  
50000/50000 [=====] 27s

Running time **without GPU**

Running time **with GPU**

With GPU, the running time is  $733/27=27.1$  times faster than the running time without GPU!!!



# TECHNOLOGICAL ADVANCEMENTS IN SOFTWARE

---



# WHY DEEP LEARNING IS SO POPULAR

TECHPRO  
EDUCATION



Keras



TensorFlow

Caffe

PyTorch



## PROMINENT FIGURES OR GURUS IN THE FIELD OF DEEP LEARNING

---



# Geoffrey Hinton

## “The Godfather of deep learning”





deeplearning.ai presents  
Heroes of Deep Learning


# Ian Goodfellow

Research Scientist at Google Brain



# Fei-Fei Li

Sequoia Professor of Computer  
Science at Stanford University





# DEEP LEARNING



deeplearning.ai



Carnegie Mellon University  
Machine Learning

“

Just as electricity transformed almost everything 100 years ago, today I actually have a hard time thinking of an industry that I don't think AI will transform in the next several years.

~ Andrew Ng






1 0 0 0 1 1 0 1 0 0  
1 0 0 0 1 1 0 1 0 0  
1 0 0 0 1 1 0 1 0 0  
1 0 0 0 1 1 0 1 0 0  
1 0 0 0 1 1 0 1 0 0  
1 0 0 0 1 1 0 1 0 0

1 1 1 1 1 0 0 1 1 0  
1 1 1 1 1 0 0 1 1 0  
1 1 1 1 1 0 0 1 1 0  
1 1 1 1 1 0 0 1 1 0  
1 1 1 1 1 0 0 1 1 0  
1 1 1 1 1 0 0 1 1 0



# DEEP LEARNING HISTORY

## Popularity




## Timeline of AI Development

- **1950s-1960s:** First AI boom - the age of reasoning, prototype AI developed
- **1970s:** AI winter I
- **1980s-1990s:** Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
- **1990s:** AI winter II
- **1997:** Deep Blue beats Gary Kasparov
- **2006:** University of Toronto develops Deep Learning
- **2011:** IBM's Watson won Jeopardy
- **2016:** Go software based on Deep Learning beats world's champions




# DEEP LEARNING HISTORY




# DEEP LEARNING HISTORY



## Deep Learning Timeline



# DEEP LEARNING HISTORY



McCulloch (right) and Pitts (left)





## Frank Rosenblatt Creates Perceptron



In his paper "The Perceptron: A Perceiving and Recognizing Automaton", Rosenblatt shows the new avatar of McCulloch-Pitts neuron – 'Perceptron' that had true learning capabilities to do binary classification on its own. This inspires the revolution in research of shallow neural network for years to come, till first AI winter.



# 1969: Perceptrons can't do XOR!



1969: Minsky and Papert proved that perceptrons cannot represent non-linearly separable target functions.



Minsky & Papert




# DEEP LEARNING HISTORY



1986

## Restricted Boltzmann Machine



Paul Smolensky comes up with a variation of Boltzmann Machine where there is not intra layer connection in input and hidden layer. It is known as Restricted Boltzmann Machine (RBM). It would become popular in years to come especially for building recommender systems.

1995

## Boltzmann Machine



David H. Ackley, Geoffrey Hinton and Terrence Sejnowski create Boltzmann Machine that is a stochastic recurrent neural network. This neural network has only input layer and hidden layer but no output layer.

**1997**

## The Milestone Of LSTM



Sepp Hochreiter and Jürgen Schmidhuber publishes a milestone paper on "Long Short-Term Memory" (LSTM). It is a type of recurrent neural network architecture which will go on to revolutionize deep learning in decades to come.

2019

## Trio Win Turing Award




**Yoshua Bengio, Geoffrey Hinton, and Yann LeCun wins Turing Award 2018 for their immense contribution in advancements in area of deep learning and artificial intelligence. This is a defining moment for those who had worked relentlessly on neural networks when entire machine learning community had moved away from it in 1970s.**




# DEEP LEARNING MILESTONES

# DEEP LEARNING MILESTONES





# DEEP LEARNING MILESTONES





# DEEP LEARNING MILESTONES


1997





# DEEP LEARNING MILESTONES

2016





# DEEP LEARNING MILESTONES

2017

